

```
In [1]: import pandas as pd
```

```
In [2]: data=pd.read_csv("/home/placement/Desktop/EEE(238)/fiat500.csv")
```

```
In [3]: data.describe()
```

```
Out[3]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [4]: data=data.drop(['ID','lat','lon'],axis=1)#unwanted columns removed
```

```
In [5]: data.shape#data['model']=data['model'].map({'lounge:1','pop:2','sport:3'})
```

```
Out[5]: (1538, 6)
```

```
In [6]: data=pd.get_dummies(data)
```

```
In [7]: data.shape
```

```
Out[7]: (1538, 8)
```

In [8]: data

Out[8]:

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [9]: #predicted value we removed from dataframe
y=data['price']
x=data.drop('price',axis=1)
```

In [10]:

y

Out[10]:

```
0      8900
1      8800
2      4200
3      6000
4      5700
```

```
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900
```

Name: price, Length: 1538, dtype: int64

In [11]:

x

Out[11]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

```
In [12]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [13]: x_test.head(5)
```

```
Out[13]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport	
	481	51	3197	120000	2	0	1	0
	76	62	2101	103000	1	0	1	0
	1502	51	670	32473	1	1	0	0
	669	51	913	29000	1	1	0	0
	1409	51	762	18800	1	1	0	0

```
In [14]: x_train.shape
```

```
Out[14]: (1030, 7)
```

```
In [15]: x_train.head()
```

```
Out[15]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport	
	527	51	425	13111	1	1	0	0
	129	51	1127	21400	1	1	0	0
	602	51	2039	57039	1	0	1	0
	331	51	1155	40700	1	1	0	0
	323	51	425	16783	1	1	0	0

```
In [16]: y_train.head()
```

```
Out[16]: 527      9990
129      9500
602      7590
331      8750
323      9100
Name: price, dtype: int64
```

```
In [17]: x_test.head()
```

```
Out[17]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

```
In [18]: y_test.head()
```

```
Out[18]: 481      7900
76      7900
1502     9400
669     8500
1409     9700
Name: price, dtype: int64
```

In [19]: x_train

Out[19]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0
...
1130	51	1127	24000	1	1	0	0
1294	51	852	30000	1	1	0	0
860	51	3409	118000	1	0	1	0
1459	51	762	16700	1	1	0	0
1126	51	701	39207	1	1	0	0

1030 rows × 7 columns

In [20]: y_train

Out[20]:

527	9990
129	9500
602	7590
331	8750
323	9100
...	...
1130	10990
1294	9800
860	5500
1459	9990
1126	8900

Name: price, Length: 1030, dtype: int64

In [21]: x_test

Out[21]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0
...
291	51	701	22000	1	1	0	0
596	51	3347	85500	1	0	1	0
1489	51	366	22148	1	0	1	0
1436	51	1797	61000	1	1	0	0
575	51	366	19112	1	1	0	0

508 rows × 7 columns

In [22]: y_test

Out[22]:

481	7900
76	7900
1502	9400
669	8500
1409	9700
...	...
291	10900
596	5699
1489	9500
1436	6990
575	10900

Name: price, Length: 508, dtype: int64

```
In [23]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()#creating object of LinearRegression
reg.fit(x_train,y_train)#training and fitting LR object using training data
```

Out[23]: LinearRegression()
In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [24]: #x_test=[[1,51,1000,28800,3],[1,51,780,18800,1]]
```

```
In [25]: #above line to actual
```

```
In [26]: ypred=reg.predict(x_test)
```

```
In [27]: ypred
```

```
Out[27]: array([ 5867.6503378 ,  7133.70142341,  9866.35776216,  9723.28874535,
        10039.59101162,  9654.07582608,  9673.14563045, 10118.70728123,
        9903.85952664,  9351.55828437, 10434.34963575,  7732.26255693,
        7698.67240131,  6565.95240435,  9662.90103518, 10373.20344286,
        9599.94844451,  7699.34400418,  4941.33017994, 10455.2719478 ,
        10370.51555682, 10391.60424404,  7529.06622456,  9952.37340054,
        7006.13845729,  9000.1780961 ,  4798.36770637,  6953.10376491,
        7810.39767825,  9623.80497535,  7333.52158317,  5229.18705519,
        5398.21541073,  5157.65652129,  8948.63632836,  5666.62365159,
        9822.1231461 ,  8258.46551788,  6279.2040404 ,  8457.38443276,
        9773.86444066,  6767.04074749,  9182.99904787, 10210.05195479,
        8694.90545226, 10328.43369248,  9069.05761443,  8866.7826029 ,
        7058.39787506,  9073.33877162,  9412.68162121, 10293.69451263,
        10072.49011135,  6748.5794244 ,  9785.95841801,  9354.09969973,
        9507.9444386 , 10443.01608254,  9795.31884316,  7197.84932877,
        10108.31707235,  7009.6597206 ,  9853.90699412,  7146.87414965,
        6417.69133992,  9996.97382441,  9781.18795953,  8515.83255277,
        8456.30006203,  6499.76668237,  7768.57829985,  6832.86406122,
        8347.96113362, 10439.02404036,  7356.43463051,  8562.56562053,
        8828.78555188, 10025.82571528,  7278.77188822,  8411.45884886])
```



```
In [28]: #savedmodel=pickle.load(open(filename,'rb'))  
#x_test=[[1,75,1062,8000,1]]  
#savedmodel.predict(x_test)
```

```
In [29]: from sklearn.metrics import r2_score  
r2_score(y_test,ypred)
```

```
Out[29]: 0.8415526986865394
```

```
In [30]: from sklearn.metrics import mean_squared_error  
t=mean_squared_error(ypred,y_test)
```

```
In [31]: import math  
print(math.sqrt(t))
```

```
762.8156575420782
```

```
In [32]: #from sklearn.metrics import accuracy_score  
#accuracy_score(y_test,ypred)
```

```
In [33]: #Results=pd.DataFrame(columns=['Actual','Predicted'])
#Results['Actual']=y_test
Results=pd.DataFrame(columns=['Price','Predicted'])
Results['Price']=y_test
Results['Predicted']=ypred
#Results['Km']=x_test['Km']
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(15)
```

```
Out[33]:
```

	index	Price	Predicted	Id
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
5	1414	9900	9654.075826	5
6	1089	9900	9673.145630	6
7	1507	9950	10118.707281	7
8	970	10700	9903.859527	8
9	1198	8999	9351.558284	9
10	1088	9890	10434.349636	10
11	576	7990	7732.262557	11
12	965	7380	7698.672401	12
13	1488	6800	6565.952404	13
14	1432	8900	9662.901035	14

```
In [34]: Results['diff']=Results.apply(lambda row: row.Price - row.Predicted,axis=1)
```

In [35]:

Results

Out[35]:

	index	Price	Predicted	Id	diff
0	481	7900	5867.650338	0	2032.349662
1	76	7900	7133.701423	1	766.298577
2	1502	9400	9866.357762	2	-466.357762
3	669	8500	9723.288745	3	-1223.288745
4	1409	9700	10039.591012	4	-339.591012
...
503	291	10900	10032.665135	503	867.334865
504	596	5699	6281.536277	504	-582.536277
505	1489	9500	9986.327508	505	-486.327508
506	1436	6990	8381.517020	506	-1391.517020
507	575	10900	10371.142553	507	528.857447

508 rows × 5 columns

In []: